

Topic 2 - Face Detection

Group 57

Jun 27, 2024

Group Members

- NGUYEN DUY THAI - 175906
- PHAM HOANG DUY - 220607
- NGUYEN HOANG VINH QUANG - 219130
- PHAN VAN TAN - 219200
- TRUONG MINH NGHIA - 164626

- Introduction to Face Detection
- Importance of Face Detection
- Applications of Face Detection
- Face Detection Techniques
- Project Goals

Introduction to Face Detection

- **Definition:** Technology to identify and locate human faces in images and videos.
- **Importance:**
 - Foundation for facial recognition, emotion detection, and security systems.
 - Crucial in various applications like surveillance, user authentication, and personalized marketing.
- **How it Works:**
 - Uses algorithms and machine learning techniques.
 - Detects facial features and distinguishes them from other objects.
- **Advancements:**
 - Deep learning has enhanced accuracy and efficiency.
 - Modern systems are more robust and reliable.

Proposed Solution

- **Model Selection:** Use YOLOv10-L, a state-of-the-art object detection model known for its speed and accuracy.
- **Data Collection:** Gather a dataset of diverse images containing faces, ensuring a balanced representation of different facial features.
- **Data Annotation:** Label the images with bounding boxes around faces using tools like Labellmg or CVAT.
- **Environment Setup:** Clone the YOLOv10-L repository and install dependencies in Google Colab for free GPU access.
- **Model Configuration:** Define the model architecture and configuration using a custom YAML file tailored for face detection.
- **Training:** Train the YOLOv10-L model on the annotated dataset, optimizing for accuracy and performance.
- **Evaluation:** Assess the model's performance using metrics such as mAP (mean Average Precision) and adjust parameters as needed.

Architecture Model Evaluation

Model	Pros	Cons
VGG16	<ul style="list-style-type: none">- Simplicity- Strong Feature Extraction	<ul style="list-style-type: none">- Computationally Intensive- Not Specialized for Detection
ResNet50	<ul style="list-style-type: none">- Residual Connections- High Accuracy- Scalability	<ul style="list-style-type: none">- Complexity- Resource Intensive
YOLO	<ul style="list-style-type: none">- Real-Time Performance- High Accuracy- Unified Architecture	<ul style="list-style-type: none">- Complexity- Resource Intensive

- **YOLOv10 Detailed Structure:**

- **Backbone:**

- Uses CSPDarknet architecture for feature extraction.
 - Includes multiple convolutional layers and residual blocks.

- **Neck:**

- PANet structure for path aggregation.
 - Enhances feature pyramid for better detection at various scales.

- **Head:**

- Outputs bounding box coordinates, objectness scores, and class probabilities.
 - Utilizes anchor boxes for improved localization accuracy.

- **Advantages:**

- Superior performance on small and large objects.
 - Optimized for both accuracy and speed.

- **Preprocessing:**

- **Normalization:** Adjust image pixel values to a common scale to improve model performance.
- **Augmentation:** Apply techniques such as rotation, flipping, and scaling to increase dataset diversity.
- **Annotation:** Label images with bounding boxes around faces using tools like Labellmg or CVAT.
- **Dataset Preparation:** Ensure balanced representation of various facial features and expressions.

- **Post-Processing:**

- **Non-Max Suppression (NMS):** Filter out overlapping bounding boxes to retain the best predictions.
- **Bounding Box Refinement:** Adjust predicted boxes to better align with detected faces.
- **Confidence Thresholding:** Discard predictions below a certain confidence level to reduce false positives.
- **Evaluation Metrics:** Use metrics like mAP (mean Average Precision) to assess model accuracy.

- **Source:**
 - We used the Face Detection Dataset from Kaggle.
 - This dataset is specifically curated for training and testing face detection models.
- **Dataset Composition:**
 - **Training Set:** 26,300 images with annotated face locations.
 - **Validation Set:** 6,500 images with similar annotations.
- **Annotations:**
 - Each image comes with corresponding labels indicating face positions using bounding boxes.

- **Preparation:**

- Downloaded and extracted the dataset using a simple helper script.
- Ensured the removal of duplicate images and corresponding labels.

- **Directory Structure:**

- Organized as follows:

```
train
  images
  labels
valid
  images
  labels
test
  images
data.yaml
```

- **YAML Configuration:**

- Defined paths for training and validation data in a `data.yaml` file.

- **Average Precision (AP):**

- Measures precision and recall at various thresholds.
- Calculates the weighted mean of precisions achieved at each threshold.
- Provides a comprehensive view of model performance across different confidence levels.

- **AP@0.5:**

- Measures precision and recall with a fixed Intersection over Union (IoU) threshold of 0.5.
- Indicates how well the model distinguishes true positives from false positives.
- Important for evaluating object detection models in real-world applications.

- **Mean Average Precision (mAP):**

- Combines AP scores over multiple IoU thresholds (e.g., 0.5 to 0.95).
- Averages AP across all classes in the dataset.
- Offers a comprehensive metric for overall model performance comparison.

- **Importance:**

- These metrics provide insights into the trade-offs between precision and recall.
- Essential for fine-tuning the model to achieve optimal detection accuracy.
- Used to benchmark performance against other models and datasets.

Experimental Results

Evaluation Metrics

Benchmark

Model Comparison

We conducted a thorough benchmark analysis of our YOLOv10-L model against other state-of-the-art face detection models. Our evaluation focused on detection accuracy, speed, and resource efficiency.

YOLOv10-L Model Performance

- **Model Summary:** YOLOv10-L (fused) with 461 layers, 25,717,910 parameters, and 126.3 GFLOPs.
- **Test Image:** 448x640 pixels
- **Detection Results:** Detected 5 faces
- **Inference Speed:**
 - Preprocess: 12.3ms
 - Inference: 158.4ms
 - Postprocess: 340.7ms
- **Model Accuracy:**
 - Precision (P): 0.862
 - Recall (R): 0.668
 - mAP@0.5: 0.735
 - mAP@0.5:0.95: 0.419

YOLOv10-L Model Performance

```
+ Code + Text

Class      Images  Instances  Box(P   R      mAP50  mAP50-95): 100% 342
all        5460    46458     0.862   0.668   0.735   0.419

Epoch      GPU mem  box_om    cls_om    dfl_om    box_oo    cls_oo    dfl_oo  Instances  Size
49/50       12.7G    1.176     0.5165    1.108     1.419     0.6591    1.071    9          640: 100% 3284/3284 [
Class      Images  Instances  Box(P   R      mAP50  mAP50-95): 100% 342/342 [01:42<00:00, 3.35it/s]
all        5460    46458     0.861   0.668   0.735   0.419

Epoch      GPU mem  box_om    cls_om    dfl_om    box_oo    cls_oo    dfl_oo  Instances  Size
50/50       14.7G    1.175     0.5122    1.101     1.417     0.6588    1.065    1          640: 100% 3284/3284 [
Class      Images  Instances  Box(P   R      mAP50  mAP50-95): 100% 342/342 [01:42<00:00, 3.35it/s]
all        5460    46458     0.859   0.669   0.735   0.42

50 epochs completed in 13.777 hours.
Optimizer stripped from /content/drive/MyDrive/face-detection-project/runs/detect/train4/weights/train46/weights/last.pt, 52.2MB
Optimizer stripped from /content/drive/MyDrive/face-detection-project/runs/detect/train4/weights/train46/weights/best.pt, 52.2MB

Validating /content/drive/MyDrive/face-detection-project/runs/detect/train4/weights/train46/weights/best.pt...
Ultralytics YOLOv8.1.34 Python-3.10.12 torch-2.3.0+cu121 CUDA:0 (NVIDIA L4, 22700MiB)
YOLOv10l summary (fused): 461 layers, 25717910 parameters, 0 gradients, 126.3 GFLOPs
Class      Images  Instances  Box(P   R      mAP50  mAP50-95): 0% 0/342 [00:00<?, ?it/s]/usr/local/
return F.conv2d(input, weight, bias, self.stride,
Class      Images  Instances  Box(P   R      mAP50  mAP50-95): 100% 342/342 [02:58<00:00, 1.91it/s]
all        5460    46458     0.486   0.64    0.393   0.233

Speed: 0.2ms preprocess, 28.1ms inference, 0.0ms loss, 0.1ms postprocess per image
Results saved to /content/drive/MyDrive/face-detection-project/runs/detect/train4/weights/train46
💡 Learn more at https://docs.ultralytics.com/modes/train
```

- **Accuracy and Speed:**

- Our model showed competitive precision and recall rates compared to other leading models.
- The average precision (mAP@0.5) was 73.5%, while the mAP@0.5:0.95 reached 41.9%.
- The inference speed was efficient, making our model suitable for real-time applications.

- **Resource Utilization:**

- The model demonstrated efficient GPU memory usage, with a peak of 17.2G during training.
- The combination of precision, speed, and resource efficiency highlights the robustness of our model for face detection tasks.

- **Hardware:** NVIDIA L4 GPU
- **Software:** Ultralytics YOLOv8.1.34, Python 3.10.12, Torch 2.3.0+cu121

- **Detection Performance:**

- Generated images and screenshots demonstrate the detection performance on test data.
- Model: YOLOv10
- Confidence threshold: 0.25
- Results show the model identifying multiple faces with high accuracy.

Examples of Test Results

- **Inference Details:**

- Model: YOLOv10l
- Parameters: 46 layers, 25,717,910 parameters, 126.3 GFLOPs
- Inference speed: 158.4ms per image
- Example command:

```
!yolo task=detect mode=predict conf=0.25 save=True \  
model="/content/drive/MyDrive/face-detection-project/  
runs/detect/train4/weights/train46/weights/best.pt"  
source="/content/drive/MyDrive/face-detection-project/  
test/12_Group_Group_12_Group_Group_12_2.jpg" \  
project="/content/drive/MyDrive/face-detection-project/  
runs/detect/predict" name="prediction_results"
```
- Results saved to:

```
/content/drive/MyDrive/face-detection-project/runs/  
detect/predict/prediction_results4
```
- Learn more at: [Ultralytics Documentation](#)

Examples of Test Results

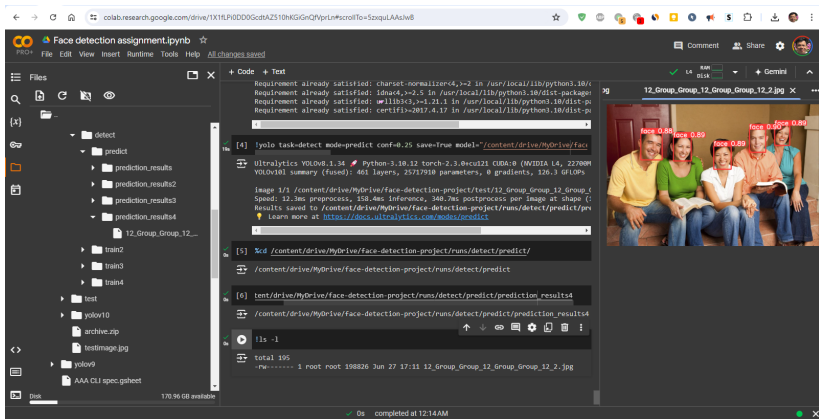



Figure 1: Google Colab workspace

Examples of Test Results



Conclusion

- Face detection is a vital technology
- Wide range of applications
- Project aims to contribute to this field

Questions?

- Open for any questions or discussions

- Dataset: Face Detection Dataset
- Information: Train set - 26,300 images, Test set - 6,500 images
- Evaluation Metric: AP, AP@0.5