

Topic 2 - Face Detection

Group 57

Jun 27, 2024

Group Members

- NGUYEN DUY THAI - 175906
- PHAM HOANG DUY - 220607
- NGUYEN HOANG VINH QUANG - 219130
- PHAN VAN TAN - 219200
- TRUONG MINH NGHIA - 164626

- Introduction to Face Detection
- Importance of Face Detection
- Applications of Face Detection
- Face Detection Techniques
- Project Goals

Introduction to Face Detection

- **Definition:** Technology to identify and locate human faces in images and videos.
- **Importance:**
 - Foundation for facial recognition, emotion detection, and security systems.
 - Crucial in various applications like surveillance, user authentication, and personalized marketing.
- **How it Works:**
 - Uses algorithms and machine learning techniques.
 - Detects facial features and distinguishes them from other objects.
- **Advancements:**
 - Deep learning has enhanced accuracy and efficiency.
 - Modern systems are more robust and reliable.

Proposed Solution

- **Model Selection:** Use YOLOv10-L, a state-of-the-art object detection model known for its speed and accuracy.
- **Data Collection:** Gather a dataset of diverse images containing faces, ensuring a balanced representation of different facial features.
- **Data Annotation:** Utilize the pre-labeled dataset downloaded from Kaggle, which includes about 32k images with clean labels in xywh (top-left x-coord, top-left y-coord, face-width, face-height) format for human face detection tasks.
- **Environment Setup:** Clone the YOLOv10-L repository and install dependencies in Google Colab for free GPU access.
- **Model Configuration:** Define the model architecture and configuration using a custom YAML file tailored for face detection.

Proposed Solution

- **Training:** Train the YOLOv10-L model on the annotated dataset, optimizing for accuracy and performance.
- **Evaluation:** Assess the model's performance using metrics such as mAP (mean Average Precision) and adjust parameters as needed.
- **Inference:** Test the trained model on new images to validate its face detection capabilities.
- **Deployment:** Export the trained model weights for use in real-world applications.

Architecture Model Evaluation

Model	Pros	Cons
VGG16	<ul style="list-style-type: none">- Simplicity- Strong Feature Extraction	<ul style="list-style-type: none">- Computationally Intensive- Not Specialized for Detection
ResNet50	<ul style="list-style-type: none">- Residual Connections- High Accuracy- Scalability	<ul style="list-style-type: none">- Complexity- Resource Intensive
YOLO	<ul style="list-style-type: none">- Real-Time Performance- High Accuracy- Unified Architecture	<ul style="list-style-type: none">- Complexity- Resource Intensive

- **YOLOv10 Detailed Structure:**

- **Backbone:**

- Uses CSPDarknet architecture for feature extraction.
 - Includes multiple convolutional layers and residual blocks.

- **Neck:**

- PANet structure for path aggregation.
 - Enhances feature pyramid for better detection at various scales.

- **Head:**

- Outputs bounding box coordinates, objectness scores, and class probabilities.
 - Utilizes anchor boxes for improved localization accuracy.

- **Advantages:**

- Superior performance on small and large objects.
 - Optimized for both accuracy and speed.

- **Preprocessing:**

- **Normalization:** Adjust image pixel values to a common scale to improve model performance.
- **Augmentation:** Apply techniques such as rotation, flipping, and scaling to increase dataset diversity.
- **Annotation:** Utilize the pre-labeled dataset downloaded from Kaggle, which includes about 32k images with clean labels.
- **Dataset Preparation:** Ensure balanced representation of various facial features and expressions.

- **Source:**
 - We used the Face Detection Dataset from Kaggle.
 - This dataset is specifically curated for training and testing face detection models.
- **Dataset Composition:**
 - **Training Set:** 26,300 images with annotated face locations.
 - **Validation Set:** 6,500 images with similar annotations.
- **Annotations:**
 - Each image comes with corresponding labels indicating face positions using bounding boxes.

- **Preparation:**

- Downloaded and extracted the dataset using a simple helper script.
- Ensured the removal of duplicate images and corresponding labels.

- **Directory Structure:**

- Organized as follows:

```
train
  images
  labels
validation
  images
  labels
test
  images
face-detect-datase.yaml
```

- **YAML Configuration:**

- Defined paths for training and validation data in a `face-detect-dataset.yaml`.
- Included class names and counts for model reference.
- **Training Data Path:**
`/content/drive/MyDrive/face-detection-project/merged/images/train`
- **Validation Data Path:**
`/content/drive/MyDrive/face-detection-project/merged/images/validation`
- **Class Names:** `['face']`
- **Number of Classes (nc):** 1

- **Average Precision (AP):**

- Measures precision and recall at various thresholds.
- Calculates the weighted mean of precisions achieved at each threshold.
- Provides a comprehensive view of model performance across different confidence levels.

- **AP@0.5:**

- Measures precision and recall with a fixed Intersection over Union (IoU) threshold of 0.5.
- Indicates how well the model distinguishes true positives from false positives.
- Important for evaluating object detection models in real-world applications.

- **Mean Average Precision (mAP):**

- Combines AP scores over multiple IoU thresholds (e.g., 0.5 to 0.95).
- Averages AP across all classes in the dataset.
- Offers a comprehensive metric for overall model performance comparison.

- **Importance:**

- These metrics provide insights into the trade-offs between precision and recall.
- Essential for fine-tuning the model to achieve optimal detection accuracy.
- Used to benchmark performance against other models and datasets.

Training Custom YOLOv10-L Detector

- **Model Summary:** YOLOv10-L with 628 layers, 25,766,870 parameters, and 127.2 GFLOPs.
- **Training:**
 - Total Epochs: 50
 - Batch Size: 8
 - Image Size: 640x640 pixels
 - Data Augmentation:
 - **Color Augmentation:** Adjustments in hue (hsv_h=0.02), saturation (hsv_s=0.8), and value (hsv_v=0.5).
 - **Geometric Transformations:** Includes rotation (degrees=5), translation (translate=0.2), scaling (scale=0.6), and shearing (shear=2).
 - **Perspective Transform:** Small adjustments with perspective (perspective=0.001).
 - **Flipping:** Both vertical flipping (flipud=0.1) and horizontal flipping (fliplr=0.6).

- **Training:**

- Data Augmentation:

- **Mosaic:** Combines four images into one (mosaic=1.0).
 - **MixUp:** Merges two images into one (mixup=0.2).
 - **Copy-Paste:** Pastes objects from one image into another (copy_paste=0.1).
 - **Auto Augment:** Uses RandAugment strategy for automatic augmentation.
 - **Erasing:** Randomly erases parts of images (erasing=0.5).
 - **Cropping:** Applies cropping with a fraction of 1.0 (crop_fraction=1.0).

- **Training Results:**

- Initial Epoch GPU Memory: 17.2G
- Final Epoch GPU Memory: 14.7G
- Total Training Time: 13.777 hours

- **Validation Results:**

- Precision (P): 0.861
- Recall (R): 0.669
- mAP@0.5: 0.735
- mAP@0.5:0.95: 0.42

- **Inference Speed:**

- Preprocess: 0.2ms per image
- Inference: 28.1ms per image
- Postprocess: 0.1ms per image

- **Model Performance:** Efficient detection and robust accuracy, suitable for real-time face detection applications

Experimental Results

- Command:

```
!yolo task=detect mode=train epochs=50 batch=12  
imgsz=640 plots=True model='/content/drive/MyDrive/  
face-detection-project/yolov10/weights/yolov10l.pt'  
data='/content/drive/MyDrive/face-detection-project/  
merged/face-detect-datase.yaml' project='/content/  
drive/MyDrive/face-detection-project/runs/detect/  
train4/weights' name='train4' augment=True  
hsv_h=0.02 hsv_s=0.8 hsv_v=0.5 degrees=5 translate=0.2  
scale=0.6 shear=2 perspective=0.001 flipud=0.1  
fliplr=0.6 mosaic=1.0 mixup=0.2 copy_paste=0.1  
auto_augment=randaugument erasing=0.5 crop_fraction=1.0
```

Experimental Results

```
+ Code + Text
```

	Class	Images	Instances	Box(P)	R	mAP50	mAP50-95)	100%	342
	all	5460	46458	0.862	0.668	0.735	0.419		

Epoch	GPU_mem	box_om	cls_om	dfl_om	box_oo	cls_oo	dfl_oo	Instances	Size
49/50	12.7G	1.176	0.5165	1.108	1.419	0.6591	1.071	9	640: 100% 3284/3284 [

Class	Images	Instances	Box(P)	R	mAP50	mAP50-95)	100%	342/342	[01:42<00:00, 3.35it/s]
all	5460	46458	0.861	0.668	0.735	0.419			

Epoch	GPU_mem	box_om	cls_om	dfl_oo	box_oo	cls_oo	dfl_oo	Instances	Size
50/50	14.7G	1.175	0.5122	1.101	1.417	0.6588	1.065	1	640: 100% 3284/3284 [

Class	Images	Instances	Box(P)	R	mAP50	mAP50-95)	100%	342/342	[01:42<00:00, 3.35it/s]
all	5460	46458	0.859	0.669	0.735	0.42			

50 epochs completed in 13.777 hours.

Optimizer stripped from /content/drive/MyDrive/face-detection-project/runs/detect/train4/weights/train46/weights/last.pt, 52.2MB

Optimizer stripped from /content/drive/MyDrive/face-detection-project/runs/detect/train4/weights/train46/weights/best.pt, 52.2MB

Validating /content/drive/MyDrive/face-detection-project/runs/detect/train4/weights/train46/weights/best.pt...

Ultralytics YOLOv8.1.34 Python-3.10.12 torch-2.3.0+cu121 CUDA:0 (NVIDIA L4, 22700MiB)

YOLOv10l summary (fused): 461 layers, 25717010 parameters, 0 gradients, 126.3 GFLOPs

Class	Images	Instances	Box(P)	R	mAP50	mAP50-95)	0%	0/342	[00:00<?, ?it/s]
all	5460	46458	0.486	0.64	0.393	0.233			

Speed: 0.2ms preprocess, 28.1ms inference, 0.0ms loss, 0.1ms postprocess per image

Results saved to /content/drive/MyDrive/face-detection-project/runs/detect/train4/weights/train46

Learn more at <https://docs.ultralytics.com/modes/train>

Figure 1: Train model

Experimental Results

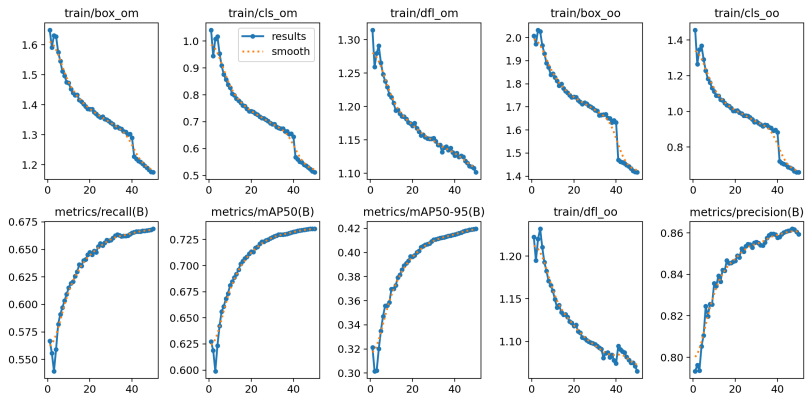


Figure 2: Train model result

- **Accuracy and Speed:**

- Our model showed competitive precision and recall rates compared to other leading models.
- The average precision (mAP@0.5) was 73.5%, while the mAP@0.5:0.95 reached 41.9%.
- The inference speed was efficient, making our model suitable for real-time applications.

- **Resource Utilization:**

- The model demonstrated efficient GPU memory usage, with a peak of 17.2G during training.
- The combination of precision, speed, and resource efficiency highlights the robustness of our model for face detection tasks.

- **Hardware:** NVIDIA L4 GPU
- **Software:** Ultralytics YOLOv8.1.34, Python 3.10.12, Torch 2.3.0+cu121
- **Detection Performance:**
 - Generated images and screenshots demonstrate the detection performance on test data.
 - Model: YOLOv10
 - Confidence threshold: 0.25
 - Results show the model identifying multiple faces with high accuracy.

Test Results: Inference on Image

- **Inference Details:**

- Model: YOLOv10l
- Parameters: 46 layers, 25,717,910 parameters, 126.3 GFLOPs
- Inference speed: 158.4ms per image
- Command:

```
!yolo task=detect mode=predict conf=0.25 save=True  
model="/content/drive/MyDrive/face-detection-project/  
runs/detect/train4/weights/train46/weights/best.pt"  
source="/content/drive/MyDrive/face-detection-project/  
test/12_Group_Group_12_Group_Group_12_2.jpg"  
project="/content/drive/MyDrive/face-detection-project/  
runs/detect/predict" name="prediction_results"
```
- Results saved to:

```
/content/drive/MyDrive/face-detection-project/runs/  
detect/predict/prediction_results4
```
- Learn more at: [Ultralytics Documentation](#)

Test Results: Inference on Image

The screenshot displays a Google Colab workspace titled "Face detection assignment.ipynb". The left sidebar shows a file explorer with a directory structure including "detect", "predict", "prediction_results", "train2", "train3", "train4", "test", "yolov10", "archive.zip", "testimage.jpg", "yolov9", and "AAA CLI spec sheet". The main code editor contains the following code:

```
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packag
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: certifi<=2017.4.17 in /usr/local/lib/python3.10/dist-p

[4] !yolo task=detect mode=predict conf=0.25 save=True model="/content/drive/MyDrive/fac
Ultralytics YOLOv8.1.34 Python-3.10.12 torch-2.3.0+cu121 CUDA:0 (NVIDIA L4, 22700M
YOLOv10l summary (fused): 461 layers, 25717910 parameters, 0 gradients, 126.3 GFLOPs

Image 1/1 /content/drive/MyDrive/Face-detection-project/test/12_Group_Group_12_Group_1
Speed: 12.3ms preprocess, 158.4ms inference, 348.7ms postprocess per image at shape (
Results saved to /content/drive/MyDrive/Face-detection-project/runs/detect/predict/pr
Learn more at https://docs.ultralytics.com/modes/cordict

[5] %cd /content/drive/MyDrive/Face-detection-project/runs/detect/predict/
/content/drive/MyDrive/Face-detection-project/runs/detect/predict

[6] !cp /content/drive/MyDrive/Face-detection-project/runs/detect/predict/prediction_results4
/content/drive/MyDrive/Face-detection-project/runs/detect/predict/prediction_result4

[7] !ls -l
total 195
-rw-rw-r-- 1 root root 198826 Jun 27 17:11 12_Group_Group_12_Group_Group_12_2.jpg
```

The terminal output shows the execution of the code, including the file listing command. The status bar at the bottom indicates "0s completed at 12:14 AM". On the right, a preview of the face detection result is shown, displaying a group of five people with bounding boxes and confidence scores (e.g., face 0.88, face 0.89, face 0.89, face 0.89, face 0.89).

Figure 3: Google Colab workspace

Test Results: Inference on Image




Figure 4: Test image 1

Test Results: Inference on Image

- **Inference Details:**

- Model: YOLOv10l
- Parameters: 461 layers, 25,717,910 parameters, 126.3 GFLOPs
- Inference speed: 155.5ms per image
- Command:

```
!yolo task=detect mode=predict conf=0.25 save=True
model="/content/drive/MyDrive/face-detection-project/
runs/detect/train4/weights/train46/weights/best.pt"
source="/content/drive/MyDrive/face-detection-project/
test/runs/AOS-group.png"
project="/content/drive/MyDrive/face-detection-project/
runs/detect/predict" name="prediction_results"
```
- Results saved to:
/content/drive/MyDrive/face-detection-project/runs/
detect/predict/prediction_results5
- Learn more at: [Ultralytics Documentation](#)

Test Results: Inference on Image



Test Results: Inference on Video

- Command:

```
!yolo task=detect mode=predict conf=0.25 save=True  
model="/content/drive/MyDrive/face-detection-project  
/runs/detect/train4/weights/train46/weights/best.pt"  
source="/content/drive/MyDrive/face-detection-project  
/test/WALK-NEW-YORK-City-USA-vlog.mp4" project="/content  
drive/MyDrive/face-detection-project/runs/detect/predict  
name="prediction_results"
```

- Log:

```
Streaming output truncated to the last 5000 lines.  
video 1/1 (frame 76799/81795) /content/drive/MyDrive/  
face-detection-project/test/WALK-NEW-YORK-City-USA-vlog  
384x640 1 face, 16.6ms
```

Test Results: Inference on Video

- Log:

```
video 1/1 (frame 76800/81795) /content/drive/MyDrive/  
face-detection-project/test/WALK-NEW-YORK-City-USA-vlog  
384x640 2 faces, 16.8ms
```

```
video 1/1 (frame 77973/81795) /content/drive/MyDrive/  
face-detection-project/test/WALK-NEW-YORK-City-USA-vlog  
384x640 (no detections), 16.7ms
```

```
...
```

```
Speed: 2.3ms preprocess, 17.3ms inference, 1.1ms  
postprocess per image at shape (1, 3, 384, 640)
```

```
Results saved to /content/drive/MyDrive/face-detection-  
project/runs/detect/predict/prediction_results6
```

- Play the inference result video:

<https://www.youtube.com/watch?v=HwbmiKk6k3I>

Test Results: Inference on Video



Figure 6: inference result video

Conclusion

- Face detection is a vital technology
- Wide range of applications
- Project aims to contribute to this field

Questions?

- Open for any questions or discussions

- Dataset: Face Detection Dataset
- Information: Train set - 26,300 images, Test set - 6,500 images
- Evaluation Metric: AP, AP@0.5